

Using Kernel Hardening Tools: seccomp

Linux sys calls

Every Linux machine will have a kernel space on the hardware, then we have a user space where we keep our apps, every app perform some calls like (read, send, close, write, execute, accept).

Seccomp purpose is to restrict Container's (pod is container) syscalls so it cannot compromise the machine.

Seccomp restricts syscalls to resources.

(calls like: clock_adjtime, mount, unmount, open....etc)

Objective is to secure/restrict the containers by what they are allowed to do?

Like AppArmor it is also configured through profiles

SecComp Profile are a list of SysCalls to allow/deny/audit.

SecComp Profiles are enabled at kubelet level config (disabled by default)

let say container run time is restricting system calls, we can enable by:

--feature-gates=SeccompDefault=true

seccomp profiles default placement:

/var/lib/kubelet/seccomp/profiles/<profile-name.json>

SECCOMP Profiles are 3 types

- Default Profile
- Audit - audit.json = all the syscalls will be logged at /var/log/syslog
- Violation- violation.json = all syscalls will be rejected.
- Fine-Grained (Custom) - fine-grained.json = I want to allow only 5 or 10 calls rest all are restricted.

SecComp profile has defaultAction and Action

SECCOMP can operate with 3 modes

Mode 0 - disabled mode

Mode 1 - strict mode

Mode 2 - filtered mode

```
{  
  "defaultAction": "SCMP_ACT_LOG"  
}
```

```
{  
  "defaultAction": "SCMP_ACT_ERRNO"  
}
```

```
{  
  "defaultAction": "SCMP_ACT_ERRNO",  
  "architectures": [  
    "SCMP_ARCH_X86_64",  
    "SCMP_ARCH_X86",  
    "SCMP_ARCH_X32"  
  ],  
  "syscalls": [  
    {  
      "names": [  
        "accept4",  
        "epoll_wait",  
        "pselect6",  
        "futex",  
        "madvise",  
        "epoll_ctl",  
        "getsockname",  
        "setsockopt",  
        "vfork",  
        "mmap",  
        "read",  
        "write",  
        "close",  
        "arch_prctl",  
        "sched_getaffinity",  
        "munmap",  
        "brk",  
        "rt_sigaction",  
        "rt_sigprocmask",  
        "sigaltstack",  
        "gettid",  
        "clone",  
        "bind",  
        "socket",  
        "openat",  
        "readlinkat",  
        "exit_group",  
        "epoll_create1",  
        "listen",  
        "rt_sigreturn",  
        "sched_yield",  
        "clock_gettime",  
        "connect",  
        "dup2",  
        "epoll_pwait",  
        "execve",  
        "exit",  
        "fcntl",  
      ]  
    }  
  ]  
}
```

```

        "getpid",
        "getuid",
        "ioctl",
        "mprotect",
        "nanosleep",
        "open",
        "poll",
        "recvfrom",
        "sendto",
        "set_tid_address",
        "setitimer",
        "writev"
    ],
    "action": "SCMP_ACT_ALLOW"
}
]
}

```

we will follow the procedure as follow after creating a cluster:

Create Profile

Copy SecComp profiles (to all nodes of cluster or the pod you want to
deploy it in)

Apply profile to Pod SecContext

Containers & System are Secure

there are three types:

Localhost

RuntimeDefault

unconfined

ApiVersion: v1

kind: Pod

metadata:

name: audit-pod

labels:

app: audit-pod

spec:

securityContext:

seccompProfile:

type: Localhost

localhostProfile: profiles/audit.json

containers:

- **name:** test-container

image: hashicorp/http-echo:1.0

args:

- "-text=just made some syscalls!"

securityContext:

allowPrivilegeEscalation: false

Make sure Privilege Escalation is false in this container as one pod can have multiple containers.

DEMO

```
kubectl get nodes
```

```
kubectl get pods
```

```
systemctl status kubelet
```

```
clear
```

```
cat /var/lib/kubelet/config.yaml
```

```
cd /var/lib/kubelet/
```

```
mkdir -p /var/lib/kubelet/seccomp/profiles
```

```
cd seccomp
```

```
ls
```

```
curl -L -o profiles/audit.json
```

```
https://k8s.io/examples/pods/security/seccomp/profiles/audit.json
```

```
curl -L -o profiles/violation.json
```

```
https://k8s.io/examples/pods/security/seccomp/profiles/violation.json
```

```
curl -L -o profiles/fine-grained.json
```

```
https://k8s.io/examples/pods/security/seccomp/profiles/fine-grained.json
```

```
ls profiles
```

Shift Terminal

```
kubectl get pod
```

```
kubectl create deployment web --image nginx --replicas 1
```

```
kubectl get pod
```

```
kubectl apply -f https://k8s.io/examples/pods/security/seccomp/ga/audit-pod.yaml
```

```
kubectl get pod -o wide
```

Shift Terminal

```
tail -f /var/log/syslog # these are the syscalls audit pod is making.
```

Shift terminal

```
kubectl get svc
```

```
kubectl expose pod audit-pod --type NodePort --port 5678
```

```
kubectl get svc
```

```
kubectl get node -o wide
```

```
curl http://192.168.56.10:30040/
```

```
# it will log every syscall
```

```
kubectl delete service audit-pod --wait
```

```
kubectl delete pod audit-pod --wait --now
```

```
#lets move to violation pod, nothing will work as it will block all the calls
```

```
kubectl apply -f https://k8s.io/examples/pods/security/seccomp/ga/violation-pod.yaml
```

```
kubectl get pod -o wide
```

```
kubectl get svc
```

```
#change terminal and you will see Violation pod
```

```
kubectl describe pod violation-pod
```

```
#backoff starting the failed container it happens because we are using violation pod which is blocking every sys call.
```

```
kubectl delete pod violation-pod --wait --now
```

```
kubectl get pod -o wide
```

```
kubectl apply -f https://k8s.io/examples/pods/security/seccomp/ga/fine-pod.yaml
```

```
kubectl get pod -o wide
```

```
kubectl get svc
```

```
kubectl expose pod fine-pod --type NodePort --port 5678
```

```
kubectl get svc
```

curl <http://192.168.56.10:31994/>

cat /var/log/syslog | grep -i fine-pod

<https://kubernetes.io/docs/tutorials/security/seccomp/>