



# **Certificate of Cloud Security Knowledge (CCSK)**

**Notes by Al Nafi**

**Domain 8**

## **Cloud Workload Security**

**Author:**

**Suaira Tariq Mahmood**

# Securing Serverless and Function as a Service

Serverless computing, also known as Function as a Service (FaaS), represents a paradigm shift in cloud computing where applications are executed as discrete functions without the need to manage underlying infrastructure. Cloud providers automatically handle provisioning, scaling, and execution, allowing developers to focus solely on writing code.

While serverless architectures offer significant advantages in scalability and cost efficiency, they introduce unique security challenges that require specialized mitigation strategies. Unlike traditional workloads, serverless functions operate in short-lived execution environments, making traditional security tools and monitoring methods less effective. Additionally, because these functions rely on cloud-native services for execution, they have complex dependencies on identity and access management (IAM), environment configurations, and third-party integrations.

Securing serverless workloads involves addressing key risks associated with function execution, implementing strong access control mechanisms, and ensuring sensitive data is securely managed. This section explores these security aspects in detail, preparing the foundation for securing event-driven applications and cloud-native microservices.

---

## 8.4.1 FaaS Security Issues

Serverless architectures present unique security challenges due to their dynamic and ephemeral nature. Unlike traditional applications running on persistent virtual machines or containers, serverless functions execute on-demand, often in stateless environments, making threat detection and mitigation more complex.

One of the primary concerns in serverless security is **insecure function invocation**. Serverless functions can be triggered by a variety of sources, including APIs, cloud storage events, database modifications, and third-party services. If these triggers are not properly secured, attackers may exploit them to execute arbitrary functions, leading to unauthorized operations. Implementing **authentication and authorization mechanisms** at the function entry points prevents unauthorized invocations.

Another significant risk is **overly permissive IAM roles** assigned to serverless functions. If a function has excessive privileges, it can access unintended cloud resources, leading to privilege escalation attacks. Adopting the **principle of least privilege (PoLP)** ensures that functions only have the minimum permissions required for execution.

**Data exposure through logging and error messages** is another critical issue. Debugging information, stack traces, and logs may contain sensitive data such as API keys, tokens, or personal identifiable information (PII). Proper **log sanitization** and **centralized monitoring** using security tools like **AWS CloudTrail, Azure Monitor, and Google Cloud Logging** help mitigate this risk.

Serverless environments are also susceptible to **event injection attacks**, where attackers manipulate input events to exploit vulnerabilities in function logic. Implementing **input validation and event sanitization** reduces the likelihood of injection-based exploits, such as SQL injection or XML external entity (XXE) attacks.

Another emerging concern is **function event data poisoning**, where attackers manipulate input data used for machine learning models, analytics, or automated decision-making. Organizations must enforce **data integrity checks** and use cryptographic verification techniques to ensure event authenticity.

Since serverless functions rely heavily on third-party dependencies, they are also at risk of **supply chain attacks**. Ensuring that only **trusted libraries** are used and regularly scanning dependencies for vulnerabilities using tools like **Snyk, Dependabot, and AWS CodeGuru** helps mitigate these risks.

By addressing these security concerns through proper function invocation controls, access management, and secure logging practices, organizations can enhance the security of their serverless workloads while maintaining the agility of event-driven architectures.

---

## 8.4.2 IAM for Serverless

Identity and Access Management (IAM) is one of the most crucial security components in serverless computing. Since serverless functions interact with various cloud services, databases, and third-party APIs, implementing granular access control policies is essential for reducing attack surfaces and preventing privilege escalation.

A common security misconfiguration in serverless deployments is **assigning overly permissive IAM roles to functions**. If a function is granted excessive privileges, it can be exploited to perform unauthorized actions, such as accessing sensitive data or modifying system configurations. Enforcing **least privilege policies (PoLP)** ensures that functions only have the necessary permissions required for their specific tasks. Cloud providers offer fine-grained IAM policies that allow security teams to define exact privileges, restricting actions at the API call level.

**Role-based access control (RBAC) and attribute-based access control (ABAC)** play a significant role in managing function permissions. Organizations should adopt **service-specific IAM roles** to isolate function privileges and prevent unauthorized cross-service interactions.

Another critical aspect of IAM for serverless is **secure authentication of API requests**. Functions exposed via API gateways should enforce authentication mechanisms such as **OAuth 2.0, JWT (JSON Web Tokens), or AWS IAM authentication**. Publicly accessible functions should be secured with **API keys, rate limiting, and access logging** to prevent abuse and unauthorized access attempts.

**IAM role chaining and temporary credentials** should be used instead of long-lived static credentials to minimize exposure risks. Cloud providers offer services like **AWS STS (Security Token Service), Google IAM Workload Identity Federation, and Azure Managed Identities**, allowing functions to obtain temporary, scoped credentials dynamically.

Multi-factor authentication (MFA) should be enforced for accessing IAM roles associated with critical serverless functions. Additionally, logging and monitoring IAM activities through **AWS CloudTrail, Azure Active Directory, and Google Cloud Audit Logs** helps detect and respond to suspicious authentication patterns.

By implementing strong IAM policies, organizations can prevent unauthorized access, minimize privilege escalation risks, and ensure secure function-to-service interactions in serverless environments.

---

### 8.4.3 Environment Variables & Secrets

Environment variables are commonly used in serverless applications to store configuration parameters, API credentials, database connection strings, and other sensitive data. While environment variables provide a convenient way to manage configurations, they also introduce security risks if not handled properly.

One of the primary concerns with environment variables is **plain-text storage of sensitive credentials**. Many developers mistakenly hardcode secrets into environment variables, which can be exposed in logs, debugging sessions, or cloud configuration metadata. Storing sensitive information in environment variables without encryption poses a significant risk of credential leakage.

A best practice for securing secrets in serverless applications is to use **dedicated secret management solutions** instead of environment variables. Cloud providers offer **AWS Secrets Manager, Azure Key Vault, and Google Secret Manager**, which provide secure, encrypted storage for API keys, database passwords, and other sensitive data. Functions can securely retrieve secrets at runtime without exposing them in plaintext.

**Access control for secrets** is another critical consideration. Only authorized functions should be permitted to access specific secrets. Implementing **fine-grained IAM policies** ensures that functions can only retrieve the credentials they require. **Role-based secret access policies** prevent unauthorized secrets exposure within multi-function applications.

To further enhance security, secrets should be **rotated regularly**. Automated secret rotation ensures that credentials remain fresh and minimizes the impact of a potential compromise. Cloud secret management services offer built-in secret rotation features, integrating with databases and authentication systems to update credentials dynamically.

**Encryption in transit and at rest** is essential for securing environment variables and secrets. Cloud-native security features such as **AWS KMS (Key Management Service), Azure Managed HSM, and Google Cloud KMS** enable organizations to encrypt sensitive data before storage and securely transmit it between services.

Developers should also **sanitize logs and debugging output** to prevent accidental exposure of sensitive information. Serverless logging platforms, including **AWS CloudWatch, Azure Monitor, and Google Cloud Logging**, should be configured to exclude secrets from log entries.

By implementing robust secret management strategies, organizations can mitigate the risks associated with storing and accessing sensitive data in serverless applications, ensuring that credentials and configurations remain secure.

---

## Conclusion

Securing serverless workloads requires a holistic approach, addressing unique security challenges such as function invocation control, access management, and sensitive data protection. By implementing strong authentication, enforcing least privilege IAM policies, and using secure secret management solutions, organizations can effectively mitigate serverless security risks.