



# **Certificate of Cloud Security Knowledge (CCSK)**

**Notes by Al Nafi**

**Domain 9**

## **Cloud Data Encryption at Rest**

**Author:**

**Zunaira Tariq Mahmood**

## 9.3 Cloud Data Encryption at Rest

This section builds upon the fundamental concepts of data protection covered in previous topics (specifically the encryption of data in transit) and addresses how to secure cloud data when it is stored. As organizations continue to adopt cloud services, ensuring that data is encrypted at rest is critical for maintaining confidentiality, meeting compliance requirements, and safeguarding against unauthorized access. The concepts here also set the stage for upcoming discussions around more specialized encryption methods and integration with broader cloud security strategies.

---

### 9.3.1 Application Level Encryption

Application Level Encryption involves encrypting data within the application or at the data source level before it is stored in any medium. This approach ensures data is protected not just within the storage layer but also throughout its lifecycle within the application environment. It requires careful coordination between developers, security teams, and the underlying infrastructure.

#### 9.3.1.1 File/API Encryption

File/API Encryption addresses the encryption of discrete files and data transfers via application programming interfaces (APIs).

- When data is generated or processed by an application, it can be encrypted at the file level before it is transmitted to cloud storage.
- This method is especially relevant for organizations that handle sensitive documents, reports, or files that must be transferred securely.
- Examples include encrypting large media files before uploading to cloud storage services and using secure API calls with embedded encryption functions.

#### 9.3.1.2 Database Encryption

Database Encryption focuses on encrypting data within relational or NoSQL databases. It can be implemented at multiple layers (application-layer encryption, transparent data encryption, or column-level encryption). Database encryption helps:

- Protect stored records (e.g., personally identifiable information or financial data).
- Meet compliance regulations that require data encryption at rest.

- Reduce the risk of data exposure in case of unauthorized database access or misconfiguration.

### **9.3.1.3 Object Storage Encryption**

Object Storage Encryption deals with encrypting data stored as objects in services like Amazon S3, Azure Blob Storage, or Google Cloud Storage. Application-level control over how objects are encrypted allows:

- Greater customization of key management, including the ability to manage separate encryption keys for different object sets.
- Strong auditability, since encryption operations and key usage are logged at the application level.
- Support for cross-region replication while maintaining encryption integrity.

### **9.3.1.4 Volume Encryption in the Cloud**

Volume Encryption encrypts entire virtual disk volumes at the application or operating system level. This ensures that all data in a virtual machine's storage—regardless of its format or file system—is protected. It is particularly important when:

- Using infrastructure-as-a-service (IaaS) environments that provide persistent virtual disk storage.
- Protecting sensitive workloads that require full disk or block-level encryption to guard against unauthorized snapshotting or physical access to underlying hardware.

---

## **9.3.2 Cloud Data Key Management Strategies**

Key management is essential to encryption at rest because encryption effectiveness depends on how keys are generated, stored, rotated, and revoked. Effective strategies balance security with operational feasibility, ensuring that keys remain under appropriate control while supporting continuous service availability.

### **9.3.2.1 Client-Side Encryption**

Client-Side Encryption is performed on the data before sending it to the cloud. The encryption keys remain on the client side. This approach:

- Places full responsibility for key security on the customer.
- Limits the cloud provider's visibility into the content of the data.
- Reduces risk associated with cloud provider compromise or misconfiguration because unencrypted data never leaves the client environment.

### 9.3.2.2 Server-Side Encryption

Server-Side Encryption encrypts data after it arrives at the cloud provider but before it is written to disk. It is often transparent to the end user. Benefits include:

- Simplicity in integration, as providers automatically handle key management and encryption/decryption processes.
- Built-in compliance controls, as major providers offer compliance certifications for their managed server-side encryption services.
- Potential limitations in customization, as organizations rely on the provider's encryption mechanisms.

### 9.3.2.3 Customer-Managed Encryption Keys

Customer-Managed Encryption Keys (CMEK) allow organizations to supply and manage their own keys, typically integrated with a cloud provider's Key Management Service (KMS). CMEK enables:

- Greater control over when and how keys are used.
- Clear audit trails for key usage, rotations, and retirements.
- Alignment with organizational security policies for separation of duties (e.g., a dedicated cryptographic key management team).

### 9.3.2.3 Customer-Managed Encryption Keys

(This subheading repeats, underscoring the importance of CMEK in different contexts or platforms.)

When utilizing multiple cloud environments, organizations may have distinct approaches to CMEK that reflect each provider's capabilities. This second mention highlights:

- The need for consistent multi-cloud key management policies.
- Ensuring portability of cryptographic operations and alignment with organizational compliance across different environments.

### 9.3.2.5 Custom Application Level Encryption

Custom Application Level Encryption refers to using proprietary or specialized encryption libraries within the application stack. This approach can include:

- Customized algorithms or encryption libraries beyond standard provider offerings.
- Additional security layers that integrate with specific business logic (e.g., encrypting certain fields based on user roles or dynamic rules).
- Potential complexity in maintenance, requiring in-house cryptographic expertise and rigorous testing to avoid implementation flaws.

---

### 9.3.2 Data Encryption Recommendations

These recommendations build on the strategies and tools outlined above. They help ensure that encryption at rest is aligned with organizational, regulatory, and operational requirements.

#### 9.3.2.1 Key Management Services (KMS)

Organizations should consider using robust KMS solutions, offered either by cloud providers or third-party vendors. Key points:

- KMS often integrates seamlessly with various cloud services.
- Automated key rotation policies enhance security.
- Access control and audit logging are essential to prevent unauthorized key usage.

#### 9.3.2.2 SaaS Considerations

When using Software-as-a-Service (SaaS), organizations often have limited control over how the SaaS provider handles encryption. Important considerations include:

- Evaluating whether the SaaS provider offers encryption at rest by default.
- Verifying the availability of features like bring-your-own-key (BYOK) or CMEK.
- Reviewing the SaaS provider's compliance attestations and third-party audits (e.g., SOC 2, ISO 27001).

#### 9.3.2.3 Default Encryption

Many cloud providers enable default encryption for data at rest, reducing the administrative overhead on customers. While convenient, organizations should:

- Confirm that default algorithms and key lengths satisfy policy requirements.
- Understand how keys are stored, rotated, and retired by the provider.
- Assess any potential impact on performance or data retrieval latencies.

#### 9.3.2.4 Different Keys for Services

Using distinct keys for different services (e.g., separate keys for object storage, databases, or virtual machine volumes) can help isolate breaches and compartmentalize data. This approach:

- Simplifies revocation and rotation procedures on a per-service basis.
- Minimizes the blast radius of a compromised key.
- Facilitates auditing and compliance reporting by service or data type.

#### 9.3.2.5 IAM Policies on Keys

Proper Identity and Access Management (IAM) controls around encryption keys are crucial to prevent unauthorized use or disclosure. Strong IAM policies:

- Define which roles can create, modify, or delete keys.
- Enforce strict least privilege access to key management functions.
- Provide audit trails to detect suspicious behavior or unauthorized key access.

#### 9.3.2.6 Alignment with Threat Models

Encryption strategies should be evaluated against an organization's threat model. Threat modeling ensures that:

- Chosen encryption methods defend against the most pressing risks (e.g., insider threats, cloud provider compromises, regulatory non-compliance).
- Controls are tested under scenarios relevant to the organization's industry, scale, and regulatory landscape.
- Mitigation strategies remain aligned with evolving threats and newly introduced cloud services.

---

### Case Study: Protecting Sensitive Healthcare Records in a Multi-Cloud Environment

#### Background:

A healthcare provider handles patient records subject to HIPAA compliance. They use a combination of AWS and Azure for storing patient information, with an on-premises data center for legacy systems.

#### Implementation Steps:

- The organization adopts Customer-Managed Encryption Keys (CMEK) in both AWS (using AWS KMS with CMEK) and Azure (using Azure Key Vault with customer-owned keys).
- Database encryption is enabled using both transparent data encryption (Azure SQL Database) and custom application-level encryption for particularly sensitive patient fields.
- Distinct keys are designated for object storage (e.g., medical images in AWS S3) versus database records (patient data in Azure SQL), allowing faster revocation in case of a key compromise in one environment.
- IAM policies limit key access to a dedicated security group, ensuring that database administrators cannot directly access the keys.

#### Outcome:

- The healthcare provider meets HIPAA compliance requirements for data encryption and key management.
- The risk of data exposure in the event of a misconfiguration or unauthorized access is minimized.

- Strong audit logging and segregation of duties allow quick detection and response to suspicious activity.

**Reference and Additional Case Study Links:**

- **National Institute of Standards and Technology (NIST) Special Publication 800-111 (Guide to Storage Encryption Technologies for End User Devices):**

<https://nvlpubs.nist.gov/nistpubs/Legacy/SP/nistspecialpublication800-111.pdf>

- **Cloud Security Alliance (CSA) Guidance on Key Management in Multi-Cloud:**

<https://cloudsecurityalliance.org/>

- **HIPAA Journal Case Study on Cloud Encryption Failures in Healthcare:**

<https://www.hipaajournal.com/>

These notes are designed to function independently by covering fundamental concepts, best practices, and practical considerations around cloud data encryption at rest. At the same time, they build on prior discussions regarding data security in transit and pave the way for subsequent topics that may explore advanced cryptographic techniques, multi-party computations, or compliance-specific encryption practices within the broader CCSK series.